

CSS 430

Program 3: Synchronization

Due Date: See the syllabus

1. Purpose

In this assignment, you will implement the monitors utilized by ThreadOS. While standard Java monitors are only able to wake up one (using *notify*) or all (using *notifyall()*) sleeping threads, the ThreadOS monitors (implemented in *SynchQueue.java*) allow threads to sleep and wake up on a specific condition. These monitors are used to implement two separate but key aspects of ThreadOS.

Firstly, using this *SynchQueue.java* monitor, you will implement *SysLib.join()* and *SysLib.exit()* system calls. The *SysLib.join()* system call will allow the parent thread to wait until a child thread terminates. Secondly, in the Disk IO subsystem, the monitors are used to prevent threads from busy waiting on disk read and write operations. Specifically, in the assignment you will block the current thread if it attempts to execute a disk read or write operation on a busy disk, put it into the *SynchQueue*'s FIFO list, and continue with other scheduled threads for execution. In this way, one can prevent I/O-bound threads from wasting CPU power.

2. SysLib.join() and SysLib.exit()

In Java, you can use the *join(Thread t)* method to have the calling thread wait for the termination of the thread pointed to by the argument *t*. However, if an application program forks two or more threads and it simply wants to wait for all of them to complete, waiting for a particular thread is not a good solution. In Unix/Linux, the *wait()* system call is based on this idea. It does not receive any arguments, (thus no PID to wait on), but simply waits for one of the child processes and returns a PID that has woken up the calling process.

In the first part of this lab you will use the *SyncQueue* Monitor to implement the *SysLib.join()* call in ThreadOS. Users of *ThreadOS* should not use the Java *join()* call directly but should utilize the *SysLib.join()* system call. We would like the *SysLib.join()* system call to follow similar semantics as the Unix/Linux *join()* system call. *SysLib.join()* will permit the calling thread to sleep until one of its child threads terminates by calling *SysLib.exit()*. It should return the ID of the child thread that woke up the calling thread.

As noted, you will use *SynchQueue.java* for implementing the *SysLib.join()* system call. *Kernel.java* should instantiate a new *SyncQueue* object called *waitQueue*. This *waitQueue* will use each thread ID as an independent waiting condition.

```
SyncQueue waitQueue = new SyncQueue(scheduler.getMaxThreads());
```

When *SysLib.join()* is invoked, *Kernel.java* will put the current thread to sleep and utilize the *waitQueue* to keep track of the thread. The condition used by *waitQueue* to track the thread is equal to the thread's ID. When *SysLib.exit()* is invoked, *Kernel.java* utilizes the *waitQueue* to wake up the thread waiting under the condition which is equal to the current thread's parent ID. In this way, the exiting thread (child) will notify the waiting thread (parent). See the code below for details:

```

1  case WAIT:
2      // get the current thread id (use Scheduler.getMyTcb( ) )
3      // let the current thread sleep in waitQueue under the condition
4      // = this thread id (= tcb.getTid( ) )
5      return OK; // return a child thread id who woke me up
6  case EXIT:
7      // get the current thread's parent id (= tcb.getPid( ) )
8      // search waitQueue for and wakes up the thread under the condition
9      // = the current thread's parent id
10     return OK;
```

Please note that one key feature of *SysLib.join()* is to return to the calling thread the ID of the child thread that has woken it up.

What type of underlying support do we need to block and wake-up the threads for *SysLib.join()* and *SysLib.exit()*? In Java, all objects come equipped with monitors. A thread can put itself to sleep inside the monitor by obtaining the monitor with the *synchronized* keyword and calling the *wait()* method. A thread can be signaled and woken-up by obtaining the monitor and calling the *notify()* method. However, Java monitors do not allow for different conditions to be signalled. Therefore, we need something a bit more for the ThreadOS situation.

In order to wake up a thread waiting for a specific condition, we want to implement a more generalized monitor. We will call this monitor, *SyncQueue*. You can use the java monitor support (*wait()* / *notify()*) to implement this class. It is also suggested that you create a class *QueueNode.java* which is utilized by *SyncQueue*. There will be one *QueueNode* object per condition supported by the *SyncQueue*. The *SyncQueue* class should provide the following methods:

Private/Public	Methods/Data	Descriptions
private	<i>QueueNode[] queue</i>	maintains an array of <i>QueueNode</i> objects, each representing a different condition and enqueueing all threads that wait for this condition. You have to implement your own <i>QueueNode.java</i> . The size of the queue array should be given through a constructor whose spec is given below.
public	<i>SyncQueue()</i> , <i>SyncQueue(int condMax)</i>	are constructors that create a queue and allow threads to wait for a default condition number (=10) or a <i>condMax</i> number of condition/event types.
public	<i>enqueueAndSleep(int condition)</i>	enqueues the calling thread into the queue and waits until a given <i>condition</i> is satisfied. It returns the ID of a child thread that has woken the calling thread.
public	<i>dequeueAndWakeup(int condition)</i> , <i>dequeueAndWakeup(int condition, int tid)</i>	dequeues and wakes up a thread waiting for a given <i>condition</i> . If there are two or more threads waiting for the same <i>condition</i> , only one thread is dequeued and resumed. The FCFS (first-come-first-service) order does not matter. This function can receive the calling thread's ID, (<i>tid</i>) as the 2nd argument. This <i>tid</i> will be passed to the thread that has been woken up from <i>enqueueAndSleep</i> . If no 2nd argument is given, you may regard <i>tid</i> as 0.

3. Disk.java

Disk.java simulates a slow disk device composed of 1000 blocks, each containing 512 bytes. Those blocks are divided into 100 tracks, each of which thus includes 10 blocks. This disk provides the following commands (all of which return a Boolean value):

Operations	Behaviors
<i>read(int blockId, byte buffer[])</i>	reads data into <i>buffer[]</i> from the disk block specified by <i>blockId</i>
<i>write(int blockId, byte buffer[])</i>	writes the <i>buffer[]</i> data to the disk block specified by <i>blockId</i>
<i>sync()</i>	writes back all logical disk data to the file <i>DISK</i>
<i>testAndResetReady()</i>	tests the disk status to see if the current disk command such as <i>read</i> , <i>write</i> , or <i>sync</i> has been completed or not, and resets it if the command has been completed.
<i>testReady()</i>	tests the disk status to see if the current disk command such as <i>read</i> , <i>write</i> , or <i>sync</i> has been completed or not. (The disk status will however not be reset.)

The DISK is created when first booting up ThreadOS (java Boot). Upon this first invocation, *Disk.java* checks if your current directory includes the file named *DISK*. If such a file exists, *Disk.java* reads data from this file into its private array that simulates disk blocks. Otherwise, *Disk.java* creates a new *DISK* file in your current directory. Every time it accepts a *sync* command, *Disk.java* writes all its internal disk blocks into the *DISK* file. When *ThreadOS Kernel* is shut down, it automatically issues a *sync* command to *Disk.java*, so that *Disk.java* will retrieve data from the *DISK* file upon the next invocation. You can always start *ThreadOS* with a clean disk by removing the *DISK* file from the directory (*rm DISK*).

The disk commands *read*, *write*, and *sync* return a Boolean value indicating if *Disk.java* has accepted the command or not. The command may not be accepted if another *read*, *write*, or *sync* is in flight. In this case *false* will be returned

and the command will need to be re-issued until the disk accepts the command and returns *true*. Note that when a *true* is returned the *read*, *write*, or *sync* has been accepted and is being executed. It is not necessarily complete. Disk IO is expensive and ThreadOS simulates the time it takes for the Disk to complete the operation and write or read from the *buffer[]*.

When the *read* and the *write* have completed and the data has been read or written from the *buffer[]* the disk completion status will be set to *true*. Similarly, the *sync* command does not guarantee that all disk blocks have been written back to the *DISK* file when it returns; one has to check the disk completion status to determine when this is ready as well. The Disk provides two commands to check the completion status: *testAndResetReady()* and *testReady()*. When the *read*, *write*, or *sync* have been completed by the disk the commands will return *true*, otherwise *false*. The *testAndResetReady()* will reset the completion status to *false* if the status had been *true*. The *testReady()* call simply tests the current status but does not reset it.

The following code shows one way to do a clean disk *read* operation. Notice that *testAndResetReady()* command is used to set the disk status so another operation can be performed.

```

1    Disk disk = new Disk( 1000 );
2
3    // a disk read operation:
4    while ( disk.read( blockId, buffer ) == false )
5        ; // busy wait
6    while ( disk.testAndResetReady( ) == false )
7        ; // another busy wait
8    // now you can access data in buffer

```

DiskIO_Snippet.java hosted with ♥ by GitHub

[view raw](#)

All disk operations are initiated through system calls from user threads. There are handled like all system calls by the ThreadOS Kernel. In *Kernel.java* the code which executes the system call forwards the disk operations to the disk device, (i.e., *Disk.java*). You can see the above read pattern used in *Kernel.java*. The following code is a portion of *Kernel.java* related to disk operations:

```

1  import java.util.*;
2  import java.lang.reflect.*;
3  import java.io.*;
4
5  public class Kernel {
6      // Interrupt requests
7      public final static int INTERRUPT_SOFTWARE = 1; // System calls
8      public final static int INTERRUPT_DISK    = 2; // Disk interrupts
9      public final static int INTERRUPT_IO      = 3; // Other I/O interrupts
10
11     // System calls
12     public final static int BOOT    = 0; // SysLib.boot( )
13     ...
14     public final static int RAWREAD = 5; // SysLib.rawread(int blk, byte b[])
15     public final static int RAWWRITE= 6; // SysLib.rawwrite(int blk, byte b[])
16     public final static int SYNC    = 7; // SysLib.sync( )
17
18     // Return values
19     public final static int OK = 0;
20     public final static int ERROR = -1;
21
22     // System thread references
23     private static Scheduler scheduler;
24     private static Disk disk;
25

```

```

26 // The heart of Kernel
27 public static int interrupt( int irq, int cmd, int param, Object args ) {
28     TCB myTcb;
29     switch( irq ) {
30         case INTERRUPT_SOFTWARE: // System calls
31             switch( cmd ) {
32                 case BOOT:
33                     // instantiate and start a scheduler
34                     scheduler = new Scheduler( );
35                     scheduler.start( );
36
37                     // instantiate and start a disk
38                     disk = new Disk( 100 );
39                     disk.start( );
40                     return OK;
41
42                     ...
43
44                 case RAWREAD: // read a block of data from disk
45                     while ( disk.read( param, ( byte[] )args ) == false )
46                         ; // busy wait
47                     while ( disk.testAndResetReady( ) == false )
48                         ; // another busy wait
49                     return OK;
50
51                 case RAWWRITE: // write a block of data to disk
52                     while ( disk.write( param, ( byte[] )args ) == false )
53                         ; // busy wait
54                     while ( disk.testAndResetReady( ) == false )
55                         ; // another busy wait
56                     return OK;
57
58                 case SYNC: // synchronize disk data to a real file
59                     while ( disk.sync( ) == false )
60                         ; // busy wait
61                     while ( disk.testAndResetReady( ) == false )
62                         ; // another busy wait
63                     return OK;
64             }
65             return ERROR;
66
67         case INTERRUPT_DISK: // Disk interrupts
68
69         case INTERRUPT_IO: // other I/O interrupts (not implemented)
70
71
72     }
73     return OK;
74 }
75 }

```

KernelSnippet.java hosted with ❤ by GitHub

[view raw](#)

The above code has two severe performance problems. First a user thread must repeat requesting a disk call until the disk accepts its request. Second, the user thread needs to repeatedly check if its request has been served. Each of these operations are done in a spin loop. While functionally correct, these spin loops will cause the user thread to waste CPU until either relinquishing CPU upon a context switch or receiving a response from the disk. In order to avoid this

performance penalty, you will utilize a monitor so that the thread can be enqueued and wait until the appropriate disk operation has occurred.

With the *SyncQueue.java*, we can now code more efficient disk operations.

```

Disk disk = new Disk( 1000 );
SyncQueue ioQueue = new SyncQueue( );
...
...
// a disk read operation:
while ( disk.read( blockId, buffer ) == false )
{
    ioQueue.enqueueAndSleep( 1 );    // relinquish CPU to another ready
}
// now check to ensure disk is not busy
while ( disk.testAndResetReady( ) == false )
{
    ioQueue.enqueueAndSleep( 2 );    // relinquish CPU to another rea
}
// now you can access data in buffer

```

In the above example, the condition 1 stands for that a thread is waiting for the disk to accept a request, whereas the condition 2 stands for that a thread is waiting for the disk to complete a service, (i.e., waiting for the *buffer[]* array to be read or written). The remaining problem is who will wake up a thread sleeping on this *ioQueue* under the condition 1 or 2. It is the *ThreadOS Kernel* that will receive an interrupt from the disk device. Therefore the kernel can wake up a waiting thread. You can consult the *Kernel.java* file to see how this is done.

4. Statement of Work

Part 1: Implementing *SysLib.join()* and *SysLib.exit()*

Design and code *SyncQueue.java* following the above specification. Modify the code of the *WAIT* and the *EXIT* case in *Kernel.java* using *SyncQueue.java*, so that threads can wait for one of their child threads to be terminated. Note that *SyncQueue.java* should be instantiated as *waitQueue* upon a boot, (i.e., in the *BOOT* case), as shown below:

```

1  public class Kernel
2  {
3      private static SyncQueue waitQueue;
4      ...
5
6      public static int interrupt( int irq, int cmd, int param, Object args ) {
7          TCB myTcb;
8          switch( irq ) {
9              case INTERRUPT_SOFTWARE: // System calls
10                 switch( cmd ) {
11                     case BOOT:
12                         ...
13                         waitQueue = new SyncQueue( scheduler.getMaxThreads( ) );
14                         ...
15                 }

```

SyncQueue-Wait.java hosted with ♥ by GitHub

[view raw](#)

Compile your implementations of *SyncQueue.java* and *Kernel.java*, and run *Test2.java* from the *Shell.class* to confirm:

1. *Test2.java* waits for the termination of all its five child threads, (i.e., the *TestThread2.java* threads).

2. *Shell.java* waits for the termination of *Test2.java*. *Shell.java* should not display its prompt until *Test2.java* has completed its execution.
3. *Loader.java* waits for the termination of *Shell.java*. *Loader.java* should not display its prompt (-->) until you type *exit* from the *Shell* prompt.

Use the ThreadOS-original version of *Shell.class*.

Part 2: Implementing Asynchronous Disk I/Os*

Before going onto Part 2, save your *Kernel.java* as *Kernel.old*. Thereafter, modify *Kernel.java* to use *SyncQueue.java* in the three case statements: *RAWREAD*, *RAWWRITE*, and *SYNC*, so that disk accesses will no longer cause spin loops. Note that *SyncQueue.java* should be instantiated as *ioQueue* upon a boot, (i.e., in the *BOOT* case).

```

1  public class Kernel
2  {
3      private static SyncQueue ioQueue;
4      ...
5
6      public static int interrupt( int irq, int cmd, int param, Object args ) {
7          TCB myTcb;
8          switch( irq ) {
9              case INTERRUPT_SOFTWARE: // System calls
10                 switch( cmd ) {
11                     case BOOT:
12                         ...
13                         ioQueue = new SyncQueue( );
14                     ...
15                 }
16             }
17         }
18     }

```

SyncQueue-IO.java hosted with ♥ by GitHub

[view raw](#)

Write a user-level test thread called *Test3.java* which spawns and waits for the completion of **X pairs** of threads (where **X = 1 ~ 4**), one conducting only numerical computation and the other reading/writing many blocks randomly across the disk. Those two types of threads may be written in *TestThread3a.java* and *TestThread3b.java* separately or written in *TestThread3.java* that receives an argument and conducts computation or disk accesses according to the argument. *Test3.java* should measure the time elapsed from the spawn to the termination of its child threads.

Measure the running time of *Test3.java* when running it on your *ThreadOS*. Then, replace *Kernel.java* with the older version, *Kernel.old* that does not permit threads to sleep on disk operations. Compile and run this older version to see the running time of *Test3.java*. Did you see the performance improved by your newer version? Discuss the results in your report.

To verify your own test program, you may run the professor's *Test3.class* that receives one integer, (i.e., X) and spawns X pairs of computation and disk intensive threads. **Since UW1-320 machines include multi-cores, you need to increase X up to 3 or 4 for observing the clear difference between interruptible and busy-wait I/Os.**

```

[css430@uw1-320-20 ThreadOS]$ java Boot
threadOS ver 1.0:
Type ? for help
threadOS: a new thread (thread=Thread[Thread-3,2,main] tid=0 pid=-1)
-->l Test3 3
l Test3 3
...
elapsed time = 85162 msec.
-->q

```

5. What to Turn In

- Softcopy:

1. Part 1:

- Kernel.old (which you have modified for Part 1)
- SyncQueue.java
- Any other java programs you coded necessary to implement SyncQueue.java (QueueNode.java or whatever you named)
- Result when running Test2.class from Shell.class.
- Specifications and descriptions about your java programs

2. Part 2:

- Kernel.java (which you have modified for Part 2)
- Test3.java as well as TestThread3a.java / TestThread3b.java if created.
- Performance result when running Test3.java on your Kernel.old (implemented in Part 1)
- Performance result when running Test3.java on your Kernel.java (implemented in Part 2)
- Specifications and descriptions about your java program

3. Report:

- Discussions about performance results you got for part 2.

[gradeguide3.txt](#) is the grade guide for the assignment 3.

6. FAQ

This website could answer your questions. Please click [here](#) before emailing the professor :-).

* RACE CONDITION:

There is a race condition in the Disk class which is found in the initial ThreadOS directory. This race may be encountered in your testing for part 2 of this assignment. No other labs are susceptible to this race. I will go over the details of the race in class but it is not required that you understand it for this this lab.

If you hit this race condition, ThreadOS will hang and a simple re-cycle will get it started again. The race is infrequent enough that you can complete your assignment as stated above with current implementations. However, if you want to avoid the race then you can download a new version the Disk class which has this fixed. In order to utilize this new Disk class you will also need a new SysLib and Kernel class. All three .java files and .class will be posted on Canvas under the Files section.

As with all distributed Kernel.java's the code does not map directly to the one produced in the Kernel.class as the student is required to write some of this code.