

CSS 430

Program 4: Paging

Due Date: See the syllabus

1. Purpose

This assignment focus on page replacement mechanisms and the performance improvements achieved by implementing a buffer cache that stores frequently-accessed disk blocks in memory. In this assignment you will implement the *enhanced second-chance* algorithm as described in the OS textbook - Chapter 9.4. You will then measure its performance when running various test cases, and consider the merits and demerits of the implementation.

2. Disk Caching

Caching data from slower external main memory to the faster internal CPU memory takes advantage of both spatial and temporal *locality* of data reference. User programs are likely to access the same data or the memory locations very close to the data previously accessed. This is called spatial locality and is a key reason why an operating system reads a block or page of data from the disk rather than reading just the few bytes of data the program has accessed. User programs also tend to access the same data again in a very short period of time which in turn means that the least-recently used blocks or pages are unlikely to be used and could be *victims* for page replacement.

To accelerate disk access in *ThreadOS*, you will implement a cache that stores frequently-accessed disk blocks into main memory. Therefore, subsequent access to the same disk block can quickly read the data cached in the memory. However, when the disk cache is full and another block needs to be cached, the OS must select a victim to replace. You will employ the *enhanced second-chance algorithm* to choose the block to replace. If the block to be replaced has been updated while in the cache it must be written back to disk; otherwise, it can just be overwritten with the new block. To maintain htis data consistency between disk and cached blocks, each cached block must have the following entry information:

Entry items	descriptions
block frame number	contains the disk block number of cached data. It should be set to -1 if this entry does not have valid block information.
reference bit	is set to 1 or true whenever this block is accessed. Reset it to 0 or false by the second-chance algorithm when searching a next victim.
dirty bit	is set to 1 or true whenever this block is written. Reset it to 0 or false when this block is written back to the disk.

The following instructions summarize how to read, write, cache and replace a block of data:

1. Block-read algorithm

Scan all the cache entries (sequential search is good enough for this assignment). If the corresponding entry has been found in the cache, read the contents from the cache. Otherwise, find a free block entry in the cache and read the data from the disk to this cache block. If the cache is full and there is not a free cache block, find a victim using the *enhanced second-chance algoirthm*. If this victim is dirty, you must first write back its contents to the disk and thereafter read new data from the disk into this cache block. Don't forget to set the reference bit.

2. Block-write algorithm

Scan the cache for the appropriate block (again, sequential search is adequate). If the corresponding entry has been found in the cache, write new data to this cache block. Otherwise, find a free block entry in the cache and write the data to this cache block. You do not have to write this data through to the disk device. Mark its cache entry as dirty. If you cannot find any free cache block, then find a victim using the *enhanced second-chance algorithm*. If this victim is dirty, you must first write back its contents to the disk and thereafter write new data into this cache block. Don't forget to set the reference bit.

3. Block-replacement algorithm

Follow the *enhanced second-chance algorithm* described in the texbook pages 415-416.

3. Statement of Work

Part 1: Implementation

Get a fresh copy of *ThreadOS* from the normal location (*/usr/apps/CSS430/ThreadOS/*). Copy *kernel_org.java* to *Kernel.java*. ***kernel_org.java* will be made available for the program 4 assignment date.** You can find it a copy of *kernel_org.java* on canvas in the files section.

Design a disk cache based on the enhanced second-chance algorithm and implement it in *Cache.java*. Its specification is:

methods	descriptions
Cache(int blockSize, int cacheBlocks)	The constructor: allocates a <i>cacheBlocks</i> number of cache blocks, each containing <i>blockSize</i> -byte data, on memory
boolean read(int blockId, byte buffer[])	reads into the <i>buffer[]</i> array the cache block specified by <i>blockId</i> from the disk cache if it is in cache, otherwise reads the corresponding disk block from the disk device. Upon an error, it should return false, otherwise return true.
boolean write(int blockId, byte buffer[])	writes the <i>buffer[]</i> array contents to the cache block specified by <i>blockId</i> from the disk cache if it is in cache, otherwise finds a free cache block and writes the <i>buffer[]</i> contents on it. No write through. Upon an error, it should return false, otherwise return true.
void sync() and void flush()	writes back all dirty blocks to <i>Disk.java</i> and thereafter forces <i>Disk.java</i> to write back all contents to the <i>DISK</i> file. The <i>sync()</i> method still maintains clean block copies in <i>Cache.java</i> , while the <i>flush()</i> method invalidates all cached blocks. The former method must be called when shutting down <i>ThreadOS</i> . On the other hand, the later method should be called when you keep running a different test case without receiving any caching effects incurred by the previous test.

Kernel_org.java uses your *Cache.java* in its *interrupt()* method, so that you don't have to modify the kernel.

1. **BOOT:** instantiates a *Cache.java* object, (say *cache*) with 10 cache blocks.
2. **CREAD:** is called from the *SysLib.cread()* library call. It invokes *cache.read()*.
3. **CWRITE:** is called from the *SysLib.cwrite()* library call. It invokes *cache.write()*.
4. **CSYNC:** is called from the *SysLib.csync()* library call. It invokes *cache.sync()*.
5. **CFLUSH:** is called from the *SysLib.flush()* library call. It invokes *cache.flush()*.

Part 2: Performance Test

Write a user-level test program called *Test4.java* that conducts disk validity tests and measures the performance. *Test4.java* should perform the following four tests:

1. **Random accesses:**
read and write many blocks randomly across the disk. Verify the correctness of your disk cache.
2. **Localized accesses:**
read and write a small selection of blocks many times to get a high ratio of cache hits.
3. **Mixed accesses:**
90% of the total disk operations should be localized accesses and 10% should be random accesses.
4. **Adversary accesses:**
generate disk accesses that do not make good use of the disk cache at all.

Implement all those test items in *Test4.java*. This java program should receive the following two arguments and perform a different test according to a combination of those arguments.

1. **The 1st argument:** directs that a test will use the disk cache or not. *Test4.java* uses *SysLib.rawread*, *SysLib.rawwrite*, and *SysLib.sync* system calls if the argument specifies no use of disk cache (that is, **-disabled**). Otherwise, **-enabled** specifies using *SysLib.cread*, *SysLib.cwrite*, and *SysLib.csync* system calls.
2. **The 2nd argument:** directs one of the above four performance tests.

Example:

```
$ java Boot
threadOS ver 1.0:
Type ? for help
```

```
thread0S: a new thread (thread=Thread[Thread-3,2,main] tid=0 pid=-1)
-->l Test4 [enabled|disabled] [1-5]
```

For each test, your *Test4.java* should display the average read time and average write time. Compare the performance and consider the results when you run each of the above four test cases with and without the disk cache.

4. What to Turn In

Softcopy (file uploads according to your Canvas Assignment):

1. Part 1:
 - Cache.java
2. Part 2:
 - Test4.java
 - Performance results that must test:
 - random accesses with cache disabled
 - random accesses with cache enabled
 - localized accesses with cache disabled
 - localized accesses with cache enabled
 - mixed accesses with cache disabled
 - mixed accesses with cache enabled
 - adversary accesses with cache disabled
 - adversary accesses with cache enabled
3. Report:
 - Specification/design explanation on Cache.java
 - Performance consideration on random accesses
 - Performance consideration on localized accesses
 - Performance consideration on mixed accesses
 - Performance consideration on adversary accesses

[gradeguide4.txt](#) is the grade guide for the assignment 4.

5. FAQ

This website could answer your questions. Please click [here](#) before emailing the professor :-).